

# [blockSQP\_2 user's manual]

Based on blockSQP\_2 manual by Dennis Janka

Reinhold Wittmann

October 28, 2025

## Contents

<b>1. Introduction</b>	<b>3</b>
<b>2. Building blockSQP_2</b>	<b>3</b>
2.1. Building the C++ package . . . . .	3
<b>3. Setting up a problem</b>	<b>4</b>
3.1. Sparsity format . . . . .	5
3.2. The C++ interface . . . . .	6
3.3. The Python interface . . . . .	12
3.4. The Julia interface . . . . .	14
<b>4. Options and parameters</b>	<b>14</b>
4.1. Basic options . . . . .	15
4.2. qpsol_options . . . . .	15
4.3. Common algorithmic options . . . . .	16
4.4. Convexification strategy . . . . .	16
4.5. Output settings . . . . .	16
4.6. Sizing strategies . . . . .	17
4.7. Filter line search . . . . .	17
4.8. Advanced termination and line search heuristics . . . . .	18
<b>5. Output</b>	<b>18</b>
<b>6. Extending blockSQP_2</b>	<b>18</b>
6.1. New methods . . . . .	18
6.2. Other QP solvers . . . . .	19
<b>References</b>	<b>19</b>

**A. License 20**

## 1. Introduction

`blockSQP_2` is a sequential quadratic programming method for finding local solutions of nonlinear, nonconvex optimization problems of the form

$$\min_{x \in \mathbb{R}^n} \varphi(x) \quad (1a)$$

$$\text{s.t. } b_\ell \leq \begin{bmatrix} x \\ c(x) \end{bmatrix} \leq b_u. \quad (1b)$$

It is particularly suited for—but not limited to—problems whose Hessian matrix has block-diagonal structure such as problems arising from direct multiple shooting parameterizations of optimal control or optimum experimental design problems.

`blockSQP_2` has been developed around the quadratic programming solver qpOASES [1] to solve the quadratic subproblems. Gradients of the objective and the constraint functions must be supplied by the user. The constraint Jacobian may be given in sparse or dense format. Second derivatives may be provided or approximated by a combination of SR1 and BFGS updates. Global convergence is promoted by the filter line search of Waechter and Biegler [4, 5] that can also handle indefinite Hessian approximations.

The method is described in detail in [2, Chapters 6–8]. These chapters are largely self-contained. The notation used throughout this manual is the same as in [2]. A publication [3] is currently under review.

`blockSQP_2` is published under the very permissive zlib free software license which should allow you to use the software wherever you need. The full license text can be found at the end of this document.

## 2. Building `blockSQP_2`

`blockSQP_2` currently consists of three parts.

- The C++ source files including the C++ interface.
- The Python interface.
- The Julia interface.

The Python and Julia interfaces are built separately from the C++ interface and depend on it.

### 2.1. Building the C++ package

`blockSQP_2` is built from the source files in the `src` directory. The following dependencies need to be linked.

- BLAS, LAPACK (e.g. OpenBLAS <https://github.com/OpenMathLib/OpenBLAS>)

- a QP solver. Currently, only qpOASES handles nonconvex QPs as needed by `blockSQP_2`.
- a sparse linear solver for qpOASES. Default is MUMPS <https://mumps-solver.org/index.php>. The proprietary MA57 from HSL offers higher performance.

The preprocessor flag “QPSOLVER\_QPOASES” needs to be set if qpOASES is used.

### **Loading MUMPS**

MUMPS is not thread safe. If parallel solution of QPs is to be used with MUMPS, workarounds need to be enabled. On Linux, a shared library “libdmumps\_c\_dyn.so” needs to be built from `blockSQP/src/ext/dmumps_c_dyn.cpp`, which is then loaded in separate namespaces via `dlopen`. On Windows, several identical shared libraries “libdmumps\_c\_dyn\_0.dll, libdmumps\_c\_dyn\_1.dll, ...” for loading via `LoadLibrary` need to be built. The loader code is contained in `blockSQP/src/blocksqp_load_mumps.cpp` and expects the shared libraries to be in the same folder as the executable containing it. Finally, the preprocessor flag “`DMUMPS_C_DYN`” needs to be set to activate the dynamic loading. In addition, qpOASES `SQProblemSchur` class constructor needs to be modified to accept a pointer to the MUMPS interface function and set the preprocessor flag “`SQPROBLEMSCHUR_ENABLE_PASSTHROUGH`” to. The above is already provided in the included qpOASES version and the `CMakeLists.txt`, which ensure shared libraries are copied to the correct location.

Check the `README` and `README_windows` for further build instructions.

## **3. Setting up a problem**

A nonlinear programming problem (NLP) of the form (1) is characterized by the following information that must be provided by the user:

- The number of variables,  $n$ ,
- the number of constraints,  $m$ ,
- the objective function,  $\varphi : \mathbb{R}^n \rightarrow \mathbb{R}$ ,
- the constraint function,  $c : \mathbb{R}^n \rightarrow \mathbb{R}^m$ ,
- and lower and upper bounds for the variables and constraints,  $b_\ell$  and  $b_u$ .

In addition, `blockSQP_2` requires the evaluation of the

- objective gradient,  $\nabla \varphi(x) \in \mathbb{R}^n$ , and the
- constraint Jacobian,  $\nabla c(x) \in \mathbb{R}^{m \times n}$ .

Optionally, the following can be provided for optimal performance of `blockSQP_2`:

- In the case of a block-diagonal Hessian, a partition of the variables  $x$  corresponding to the diagonal blocks,
- a function  $r$  to compute a point  $x$  where a reduced infeasibility can be expected,  $r : \mathbb{R}^n \rightarrow \mathbb{R}^n$ ,
- a partition of  $x$  into free and dependent blocks, e.g. controls and dependent stage states.  
(→ Array of `vblocks`, see 3.2.2)

### 3.1. Sparsity format

The functions `initialize` and `evaluate` can be implemented as sparse or dense, the option `sparse_mode` controls which are used.

In `blockSQP_2`, we work with the column-compressed storage format (Harwell–Boeing format). There, a sparse matrix is stored as follows:

- an array of nonzero elements `double jacNz[nnz]`, where `nnz` is the number of nonzero elements,
- an array of row indices `int jacIndRow[nnz]` for all nonzero elements, and
- an array of starting indices of the columns `int jacIndCol[nVar+1]`.

For the matrix

$$\begin{pmatrix} 1 & 0 & 7 & 3 \\ 2 & 0 & 0 & 0 \\ 0 & 5 & 0 & 3 \end{pmatrix}$$

the column-compressed format is as follows:

```

nnz=6;
jacNz [0]=1.0;
jacNz [1]=2.0;
jacNz [2]=5.0;
jacNz [3]=7.0;
jacNz [4]=3.0;
jacNz [5]=3.0;

jacIndRow [0]=0;
jacIndRow [1]=1;
jacIndRow [2]=2;
jacIndRow [3]=0;
jacIndRow [4]=0;
jacIndRow [5]=2;

jacIndCol [0]=0;
jacIndRow [1]=2;
jacIndRow [2]=3;
jacIndRow [3]=4;
jacIndRow [4]=6;
```

In `examples/example1.cc`, `initialize` and `evaluate` are implemented both sparse and dense using a generic conversion routine that converts a dense matrix (given as `Matrix`) into a sparse matrix in column-compressed format.

Note that the sparsity pattern is not allowed to change during the optimization. That means you may only omit elements of the constraint Jacobian that are *structurally* zero, i.e., that can never be nonzero regardless of the current value of `xi`. On the other hand, `jacNz` may also contain zero values from time to time, depending on the current value of `xi`.

### 3.2. The C++ interface

`blockSQP_2` is written in C++ and uses an object-oriented programming paradigm. The method itself is implemented in a class `SQPmethod`. Furthermore, `blockSQP_2` provides a basic class `Problemspec` that is used to specify an NLP of the form (1). To solve an NLP, first an instance of `Problemspec` must be passed to an instance of `SQPmethod`. Then, `SQPmethod`'s appropriate methods must be called to start the computation.

In the following, we first describe the `Problemspec` class and how to implement the mathematical entities mentioned above. Afterwards we describe the necessary methods of the `SQPmethod` class that must be called from an appropriate driver routine. Some examples where NLPs are specified using the `Problemspec` class and then passed to `blockSQP_2` via a C++ driver routine can be found in the `examples/` subdirectory.

#### 3.2.1. Class `Problemspec`

To use the class `Problemspec` to define an NLP, you must implement a derived class, say `MyProblem`, where at least the following are implemented:

1. A constructor,
2. the method `initialize`, for sparse or dense Jacobian,
3. the method `evaluate`, for sparse or dense Jacobian.

`blockSQP_2` can be used with sparse and dense variants of qpOASES. Depending on the preferred version (set by the algorithmic option `sparse_mode`, see Sec. 4.7), the constraint Jacobian must be provided in sparse or dense format by the user.

Before passing an instance of `MyProblem` to `blockSQP_2`, the following attributes must be set:

1. `int nVar` - the number of variables,
2. `int nCon` - the number of constraints (linear and nonlinear),
3. `int nnz` (optional) - number of nonzeros in sparse constraint Jacobian, not required for dense mode,
4. `Matrix lb_var, ub_var` - lower and upper bounds for variables,

5. `Matrix lb_con, ub_con` - lower and upper bounds for constraints,
6. `int nBlocks` - the number of diagonal blocks in the Hessian,
7. `int* blockIdx` - an array of dimension `nBlocks+1` with the indices of the partition of the variables that correspond to the diagonal blocks. It is required that `blockIdx[0]=0` and `blockIdx[nBlocks]=nVar`,
8. `vblock *vblocks` (optional) - an array of `vblock {int size, bool dependent}` (3.2.2) to distinguish free and dependent variables.
9. `Condenser *cond` (optional) - a Condenser instance for the problem, see sect. 3.2.2 on how to create one.

An bound is specified by setting it to  $\pm \text{inf}$  where `inf` is the value of the option `inf, std::numeric_limits<double>::infinity()` by default.

### **Matrix**

The class `Matrix` is a simple interface to facilitate maintaining dense matrices, including access to the individual elements (internally stored column major as an array of `double`). It has the constructor `Matrix(m, n=1, ldim=-1)` for constructing an  $m \times n$  `Matrix`, elements can be accessed via operator `(i, j)`.

### **Function initialize**

`initialize` is called once by `blockSQP_2` before the SQP method is started. The dense version takes the following arguments:

- `Matrix &xi`, the optimization variables
- `Matrix &lambda`, the Lagrange multipliers
- `Matrix &constrJac`, the (dense) constraint Jacobian

All variables are initialized by zero on input and should be set to the desired starting values on return. In particular, you may set parts of the Jacobian that correspond to purely linear constraints (i.e., that stay constant during optimization) here.

The sparse version of `initialize` takes the following arguments:

- `Matrix &xi`, the optimization variables
- `Matrix &lambda`, the Lagrange multipliers
- `double *jacNz`, nonzero elements of constraint Jacobian
- `int *jacIndRow`, row indices of nonzero elements
- `int *jacIndCol`, starting indices of columns

`xi` and `lambda` are initialized by zero and must be set the same as in the dense case. `jacNz`, `jacIndRow`, `jacIndCol` are the sparse Jacobian in CCS format 3.1 and already allocated. `jacIndRow` and `jacIndCol` need to be initialized here and one can initialize parts of `jacNz` that stay constant. As such, the sparsity structure cannot be changed during the iterations and therefore should include all structurally nonzero elements.

### Function `evaluate`

Similar to `initialize`, two versions of `evaluate` exist. `evaluate` is called repeatedly by `blockSQP_2` to evaluate functions and/or derivatives for different `xi` and `lambda`. The dense version takes the following arguments:

- `const Matrix &xi`, current value of the optimization variables (input)
- `const Matrix &lambda`, current value of the Lagrange multipliers (input)
- `double *objval`, pointer to objective function value (output)
- `Matrix &constr`, constraint function values (output)
- `Matrix &grad0bj`, gradient of objective (output)
- `Matrix &constrJac`, dense constraint Jacobian (output)
- `SymMatrix *hess`, (blockwise) Hessian of the Lagrangian (output)
- `int dmode`, derivative mode (input)
- `int *info`, error flag (output)

Depending on the value of `dmode`, the following must be provided by the user:

- `dmode=0`: compute function values `objval` and `constr`
- `dmode=1`: compute function values and first derivatives `grad0bj` and `constrJac`
- `dmode=2`: compute function values, first derivatives, and Hessian of the Lagrangian for the *last*<sup>1</sup> diagonal block, i.e., `hess[nBlocks-1]`
- `dmode=3`: compute function values, first derivatives, and all blocks of the Hessian of the Lagrangian, i.e., `hess[0],...,hess[nBlocks-1]`.

`dmode=2` and `dmode=3` are only relevant if the option `exact_hess_usage` is set to 1 (last block) or 2 (full Hessian). The default is 0. On return, the variable `info` must be set to 0 if the evaluation was successful and to a value other than 0 if the computation was not successful.

---

<sup>1</sup>whichSecondDerv=1 can be useful in a multiple shooting setting: There, the lower right block in the Hessian of the Lagrangian corresponds to the Hessian of the *objective*. See [2] how to exploit this for problems of nonlinear optimum experimental design.

In the sparse version of `evaluate`, the Jacobian must be evaluated in CCS format instead using the arrays `jacNz`, `jacIndRow`, and `jacIndCol` as described in 3.1. Here only `jacNz` should be set.

#### Function `reduceConstrVio`

Whenever `blockSQP_2` encounters an infeasible QP or cannot find a step length that provides sufficient reduction in the constraint violation or the objective, it resorts to a feasibility restoration phase to find a point where the constraint violation is smaller. This is usually achieved by solving an NLP to reduce some norm of the constraint violation. In `blockSQP_2`, a minimum  $\ell_2$ -norm restoration phase is implemented. The restoration phase is usually very expensive: one iteration for the minimum norm NLP is usually more expensive than one iteration for the original NLP! As an alternative, `blockSQP_2` provides the opportunity to implement a problem-specific restoration heuristic because sometimes a problem “knows” (or has a pretty good idea of) how to reduce its infeasibility<sup>2</sup>

This routine is of course highly problem-dependent. If you are not sure what to do here, just do not implement this method. Otherwise, the method just takes `xi`, the current value of the (infeasible) point as input and expects it to be mutated to a new point. In addition, it takes a flag `info` that must be set indicating if the evaluation was successful, in which case `info=0`.

#### 3.2.2. Class Condenser

The condenser class of `blockSQP_2` depends the following structs.

- `vblock`: `{int size, bool dependent}` corresponds to a block of variables and has a flag to mark dependent variables.
- `cblock`: `{int size, bool removed}` corresponds to a block of constraints and has a flag to mark used-up conditions.
- `condensing_target`: `{int n_stages, int first_free, int vblock_end, int first_cond, int cblock_end}` describes a target structure for condensing.

Variables and constraints are partitioned into `vblocks` and `cblocks`, with both being given as arrays `vblocks = vblock*`, `cblocks = cblock*` of lengths `n_vblocks`, `n_cblocks`. Multiple shooting leads to a sequence of alternating free and dependend variable blocks. The continuity conditions are a sequence of `cblocks` with sizes corresponding the the total size of one or several consecutive dependent variable blocks.

**IMPORTANT:** Right now, the continuity conditions are expected to be implemented in the form  $x^{(k+1)} - S(x^{(k)}, u^{(k)}) = 0$ . Support for conditions  $S(x^{(k)}, u^{(k)}) - x^{(k+1)}$  may be added in the future.

---

<sup>2</sup>A prominent example are dynamic optimization problems parameterized by multiple shooting: there, additional continuity constraints for the differential states are introduced that can be violated during the optimization. Whenever `blockSQP_2` calls for the restoration phase, the problem can instead try to integrate all states over the *whole* time interval and set the shooting variables such that the violation due to continuity constraints is zero. This is often enough to provide a sufficiently feasible point and the SQP iterations can continue.

Each `condensing_target` contains start and end indices of the alternating free-dependent variable blocks and `cblocks` corresponding to the respective continuity conditions.

- `n_stages` - number of shooting stages
- `first_free` - index of first free variable block that dependent blocks depend on
- `vblock_end` - index of the first variables block not dependent or depended upon
- `first_cond` - index of the constraint block that
- `cblock_end` - index of the constraint block that

The condenser class requires the `vblocks`, `cblocks` and `targets` for construction, in addition to an integer array of the Hessian block sizes.

- `vblock *vblocks, int num_vblocks`
- `cblock *cblocks, int num_cblocks`
- `int *hsizes, int n_hsizes`
- `condensing_target *targets, int num_targets`
- `int add_dep_bounds`

The final argument `add_dep_bounds` determines whether bounds of dependent variables are added as constraints to the condensed QP: 0 - not added, 1 - added, but bounds set to  $\pm\infty$ , 2 (default) - added.

The method `full_condense` can then be called to condense a QP. It takes 7 arguments and 7 equivalent return arguments. They are

- `Matrix& grad_obj` - linear term in the QP
- `Sparse_Matrix& constr_jac` - the sparse constraint Jacobian (see `blocksqp_matrix.hpp`)
- `SymMatrix* hess` - the array of Hessian blocks
- `Matrix& lb_var, ub_var, lb_con, ub_con` - variable and constraint bounds

This also stores the data necessary to recover the original solution of the specific QP given. The user is expected to solve the condensed QP, then call `recover_var_mult` with the primal-dual condensed QP solution (`Matrix xi_cond, lambda_cond`) and the original QP solution (`Matrix xi, lambda`) as return arguments.

### **Passing the condenser to the solver**

Condensing can be activated by setting the `cond` member of the `Problemspec` instance to a pointer to it.

### 3.2.3. Class SQPmethod

If you have implemented a problem using the `Problemspec` class you may solve it with `blockSQP_2` using a suitable driver program. There, you must include the header file `blocksqp_method.hpp` (and of course any other header files that you used to specify your problem). An instance of `SQPmethod` is created with a constructor that takes the following arguments:

- `Problemspec *problem`, the NLP, see above
- `SQPOptions *parameters`, an object in which all algorithmic options and parameters are stored
- `SQPstats *statistics`, an object that records certain statistics during the optimization and – if desired – outputs some of them in files in a specified directory

Instances of the classes `SQPOptions` and `SQPstats` must be created before. `SQPOptions` uses the default constructor, its fields are detailed below.

`SQPstats` takes a `char*` that describes the directory where output files are created if file output is enabled (`SQPOptions::debug_level`).

Once an `SQPmethod` has been created, the NLP can be solved by calling the following.

- `SQPmethod::init()`: Must be called before . Therein, the user-defined `initialize` method of the `Problemspec` class is called.
- `SQPmethod::run(int maxIt, int warmStart = 0)`: Run the SQP algorithm with the given options for at most `maxIt` iterations. You may call with `warmStart=1` to continue the iterations from an earlier call. In particular, the existing Hessian information is re-used. That means that

```
|| SQPmethod* method;
|| [...]
|| method->run( 2 );
```

and

```
|| SQPmethod* method;
|| [...]
|| method->run( 1 );
|| method->run( 1, 1 );
```

yield the same result. It returns an instance of the `SQPresult` enum class to indicate success (>0), max iteration reached (=0) or failure (<0). See `blocksqp_defs.hpp` for all values of `SQPresult`.

- `SQPmethod::finish()`: Should be called after the last call to `run` to make sure all output files are closed properly.

Again, we strongly recommend to study the example in `examples/example1.cpp`, where all steps are implemented for a simple NLP with block-diagonal Hessian.

### 3.2.4. Retrieving the primal-dual iterate

The primal iterate can be retrieved by calling

```
|| Matrix SQPmethod::get_xi()
```

or

```
|| void SQPmethod::get_xi(Matrix &xi_hold)
```

The Lagrange multipliers/dual iterate is retrieved the same way via the methods

```
|| Matrix SQPmethod::get_lambda()
|| void SQPmethod::get_lambda(Matrix &lambda_hold)
```

**Important:** `blockSQP_2` defines the Lagrangian as

$$L(\xi, \lambda) = f(\xi) - \lambda^T g(\xi). \quad (2)$$

Therefore, strongly active lower bounds are characterized by the associated Lagrange multiplier being  $> 0$ , strongly active upper bounds by the multiplier being  $< 0$ . Compared to optimizers defining the Lagrangian as  $L(\xi, \lambda) = f(\xi) + \lambda^T g(\xi)$ , e.g. `ipopt`[6], the multipliers `blockSQP_2` returns will have the opposite sign.

## 3.3. The Python interface

The Python interface consists of

- The compiled `py_blockSQP.so` (`py_blockSQP.cpython-... .so`).
- The `blockSQP.py` Python-side `Problemspec` class.

They both need to be imported, either directly or through a Python `__init__.py` module specification.

```
|| import numpy as np
|| import py_blockSQP
```

Like in the C++ interface, problem specification, options and statistics objects need to be created and passes to an `SQPmethod`.

```
|| opts = py_blockSQP.SQPOptions()
|| opts.opt_tol = 1e-5
|| ...
|| stats = py_blockSQP.SQPstats("./solver/output/path")
|| ...
|| prob = py_blockSQP\_\_Problemspec()
```

The following data and function attributes need to be provided.

- `nVar` [int] - number of optimization variables

- nCon [int] - number of constraints
- f [np.ndarray[np.float64](nVar, ) → float/np.float64/Cdouble] - objective function
- grad\_f [np.ndarray[np.float64]](nVar, ) → np.ndarray[np.float64](nVar, ) - objective gradient
- g [np.ndarray[np.float64]](nVar, ) → np.ndarray[np.float64](nCon, ) - constraint function
- jac\_g [np.ndarray[np.float64]](nVar, ) → np.ndarray[np.float64](nCon, nVar) - constraint Jacobian, only used and required in dense mode

If sparse derivatives should be used, the option sparse must left at True and the make\_sparse method has to be called with the following arguments

- jac\_g\_nnz [float] - the number of structural nonzero in the constraint Jacobian
- jac\_g\_row [np.ndarray[np.int32], shape = (jac\_g\_nnz,)] - row indices of constraint Jacobian in CCS format
- jac\_g\_colind [np.ndarray[np.int32], shape = (nVar + 1,)] - column start indices of constraint Jacobian in CCS format.

Finally, the bounds and Hessian block indices should be set by calling the set\_bounds method with

- lb\_var [np.ndarray[np.float64](nVar, )] - lower variable bounds
- ub\_var [np.ndarray[np.float64](nVar, )] - upper variable bounds
- lb\_con np.ndarray[np.float64](nCar, )] - lower constraint bounds
- ub\_con np.ndarray[np.float64](nCar, )] -upper constraint bounds

and set\_blockIndex with hessBlock\_index [np.ndarray[np.int32]]. A feasibility restoration heuristic can be added as the attribute `continuity_restoration` mutates an `np.ndarray` of length `nVar`.

**Important:** `prob.complete()` must be called once all data has been provided to the problem to finalize it.

`blockSQP_2` is then called as in the C++ interface

```

|| optimizer = py_blockSQP.SQPmethod(prob, opts, stats)
|| optimizer.init()
|| ret = optimizer.run(max_num_iterations)
|| optimizer.finish()
```

The optimizer returns an SQPResult enum class instance to indicate the result of the optimization, as in the C++ interface. It can be converted either to a string to yield 'SQPResult.[name]' or to an integer. The optimal primal and dual variables and be retrieved via

```

|| xi_opt = np.array(optimizer.get_xi()).reshape(-1)
|| lam_opt = np.array(optimizer.get_lambda()).reshape(-1)
```

Check the script `benchmark/run_blockSQP.py` for an example on how to invoke `py_blockSQP`.  
**Important:** Mind the sign of the returned multipliers, see Section 3.2.4.

### 3.4. The Julia interface

The Julia interface requires the compiled `blockSQP_jl.so/.dll` binary and the Julia side module `blockSQP.jl`. Here, `Problemspec` is instead named `blockSQPProblem` and most members need to be passed to the constructor.

```
using blockSQP

prob = blockSQP.blockSQPProblem(f,g, grad_f, jac_g,
                                lb_var, ub_var, lb_con, ub_con,
                                x0, lambda0, blockIdx = Int32[0, 1, 2])
blockSQP.make_sparse!(prob, Int32(nnz), jac_g_nz, jac_g_row, jac_g_colind)

QPopts = qpOASES_options(sparsityLevel = 2)
opts = blockSQPOptions(
    maxiters = 100,
    opt_tol = 1.0e-12,
    feas_tol = 1.0e-12,
    enable_linesearch = false,
    hess_approx = 1,
    fallback_approx = 2,
    print_level = 2,
    indef_delay = 1
)
stats = blockSQP.SQPsstats("./")

method = blockSQP.Solver(prob, opts, stats)

blockSQP.init!(meth)
blockSQP.run!(!, Int32(...), Int32(...))
blockSQP.finish!(meth)

x_opt = blockSQP.get_primal_solution(meth)
lam_opt = blockSQP.get_dual_solution(meth)
```

**Important:** Mind the sign of the returned multipliers, see Section 3.2.4.

An `Optimization.jl` interface is available, see `blockSQP.jl/example_rosenbrock.jl`.

## 4. Options and parameters

In this section we describe the most important options that are passed to `blockSQP_2` through the `SQPOptions` class. Read the `options.hpp` header file for an always-up-to-date list of available options.

## 4.1. Basic options

The most important options for users wanting to solve structured problems.

Name	Description/possible values	Default
<code>sparse</code>	Use dense/sparse callbacks	False
<code>exact_hess</code>	Request exact Hessian in callbacks 0: Disabled 1: Request last block 2: Request all blocks	0
<code>hess_approx</code>		1
	0: Scaled identity 1: SR1 2: damped BFGS	
<code>fallback_approx</code>	See <code>hess_approximation</code> , has to be positive definite	2
<code>block_hess</code>	enable blockwise updates 0: Full space updates 1: Partitioned updates* 2: 2 blocks - first n-1 together and last block* * requires providing <code>blockIdx</code>	1
<code>qpsol</code>	Which QP solver should be used C++: passed as <code>QPsolvers::qpOASES</code> Python/Julia: passed as "qpOASES"	qpOASES
<code>qpsol_options</code>	Options to be passed to the QP solver C++: passed as class <code>qpOASES_options</code> Python/Julia: passed as a dict	

## 4.2. `qpsol_options`

Options to be passed to the QP solver. For `qpOASES` :

Name	Description/possible values	Default
<code>sparsityLevel</code>	Which <code>qpOASES</code> solver class -1: Infer for <code>SQPOptions</code> 0: <code>qpOASES::SQProblem</code> , dense matrices 1: <code>qpOASES::SQProblem</code> , sparse matrices 2: <code>qpOASES::SQProblemSchur</code> , sparse matrices	-1
→ <b>qpOASES manual</b>		
<code>printLevel</code>	Output level of <code>qpOASES</code>	0
<code>terminationTolerance</code>		1e-10

See `options.hpp` for other QP solver's options.

### 4.3. Common algorithmic options

Name	Symbol/Meaning	Default
<code>opt_tol</code>	$\varepsilon_{\text{opt}}$	1.0e-6
<code>feas_tol</code>	$\varepsilon_{\text{feas}}$	1.0e-6
<code>eps</code>	machine precision	1.0e-16
<code>inf</code>	$\infty$	<code>double(<math>\infty</math>)</code>
<code>max_QP_it</code>	Maximum number of QP iterations per SQP iteration	5000
<code>max_QP_secs</code>	Maximum time in seconds for qpOASES per SQP iteration	10000.0

### 4.4. Convexification strategy

Regularizing indefinite SR1 or exact Hessians increases the time per SQP iteration, but tends to make the method more consistent. There are quite a few examples that rarely converge with SR1-BFGS but converge fast with SR1 Hessians regularized with scaled identities. The following options configure these convexification strategies.

Name	Description/possible values	Default
<code>conv_strategy</code>	Type of convexification strategy 0: Convex combination between Hessian and fallback 1: Add scaled identities 2: Add scaled identities for free indices* *requires providing <code>vblocks</code>	1
<code>max_conv_QPs</code>	Maximum number additional QPs per SQP iteration 1: Hessian - fallback > 1: Hessian - convexified Hessians - fallback	1
<b>Advanced options</b>		
<code>conv_tau_H</code>	Tolerance for convexified step acceptance (see paper)	2/3
<code>conv_kappa_0</code>	Initial scale for added identities (see paper [])	1/16
<code>conv_kappa_max</code>	Maximum scale for added identities	2.0
<code>par_QPs</code>	Enable parallel solution of QPs, requires special build → sect. 2.1.	False
<code>enable_QP_cancellation</code>	Terminate long-running QP threads in parallel solution of QPs.	True
<code>indef_delay</code>	Only use fallback Hessian in first # iterations	3

### 4.5. Output settings

Name	Description/possible values	Default
------	-----------------------------	---------

<code>print_level</code>	Amount of onscreen output per iteration 0: no output 1: normal output 2: verbose output	1
<code>result_print_color</code>	Control output of result at termination 0: no output 1: no output color 2: colored output	2
<code>debug_level</code>	Amount of file output per iteration 0: no debug output 1: print one line per iteration to file 2: extensive debug output to files (impairs performance)	0

#### 4.6. Sizing strategies

Oren-Luenberger [], Shanno-Phua [] and centered Oren-Luenberger [] strategies are available for sizing Hessian approximations.

Name	Description/possible values	Default
<code>initial_hess_scale</code>	Initial scaling of the Hessian approximation	1.0
<code>sizing</code>	Hessian sizing strategy 0: off 1: Shanno-Phua 2: Oren-Luenberger (OL) 3: geometric mean of 1 and 2 4: centered Oren-Luenberger (COL)	2
<code>fallback_sizing</code>	fallback Hessian sizing strategy	4
<b>Advanced options</b>		
<code>COL_eps</code>	$\text{COL } \varepsilon$	0.1
<code>COL_tau_1</code>	$\text{COL } \tau_1$	0.5
<code>COL_tau_2</code>	$\text{COL } \tau_2$	$10^4$
<code>OL_eps</code>	Oren-Luenberger sizing $\varepsilon$	$10^{-4}$

#### 4.7. Filter line search

Name	Description	Default
<code>enable_linesearch</code>	Enable filter line search	true
<code>max_linesearch_steps</code>	Max number of step reductions	10
<code>max_consec_reduced_steps</code>	Reset Hessian after # reduced steps	8
<code>max_consec_skipped_updates</code>	Reset Hessian after # skipped updates	100

<code>skip_first_linesearch</code>	Skip line search on first iteration	false
------------------------------------	-------------------------------------	-------

## 4.8. Advanced termination and line search heuristics

Name	Description	Default
<code>enable_premature_termination</code>	Allow early termination with partial success	false
<code>max_extra_steps</code>	Max steps after reaching tolerance to refine solution	0
<code>max_filter_overrides</code>	Number of times filter rules can be overridden	2

## 5. Output

When the algorithm is run, it typically produces one line of output for every iteration. The columns of the output are:

Column	Description
<code>it</code>	Number of iteration
<code>qpIt</code>	Number of QP iterations for the QP that yielded the accepted step
<code>qpIt2</code>	Number of QP iterations for the QPs whose solution was rejected
<code>obj</code>	Value of objective
<code>feas</code>	Infeasibility
<code>opt</code>	Optimality
<code> lgrd </code>	Maximum norm of Lagrangian gradient
<code> stp </code>	Maximum norm of step in primal variables
<code> lstp </code>	Maximum norm of step in dual variables
<code>alpha</code>	Steplength
<code>nSOCS</code>	Number of second-order correction steps
<code>sk</code>	Number of Hessian blocks where the update has been skipped
<code>da</code>	Number of Hessian blocks where the update has been damped
<code>sca</code>	Value of sizing factor, averaged over all blocks
<code>QPr</code>	Number of QPs whose solution was rejected

## 6. Extending blockSQP\_2

### 6.1. New methods

Due to blockSQPs object oriented design, it is possible to extend the solver by subclassing `SQPmethod`. The `bound_correction_method` is such an example. It implements the minimalist bound correction strategy discussed in the `blockSQP_2` paper. Further extensions can be made by adding function calls to the main loop and adding more fields to the `SQPOptions` class. `blockSQP_2` is actively being developed and is continuously made more fine-grained to enhance modularity and extensibility.

## 6.2. Other QP solvers

QP solvers are added to `blockSQP_2` by subclassing the abstract `QPsolver` class and filling out its methods. The code should be included based on a preprocessor definition named `QPSOLVER_[name of QP solver]` such as the existing `QPSOLVER_qpOASES`. In addition, the classes defined in `options.h` need to be extended to account for the newly available QP solver. The name of the solver should be added to the enum class `QPsolvers` and a `QPsolver_options` subclass should be created to enable passing options to the QP solver. The interfaces to other languages need to be updated consistent with the handling of existing QP solvers. This way, experimental interfaces to `gurobi` and `QPALM` are already implemented.

## References

- [1] Hans Joachim Ferreau, Christian Kirches, Andreas Potschka, Hans Georg Bock, and Moritz Diehl. qpOASES: A parametric active-set algorithm for quadratic programming. *Mathematical Programming Computation*, pages 1–37, 2014.
- [2] Dennis Janka. *Sequential quadratic programming with indefinite Hessian approximations for nonlinear optimum experimental design for parameter estimation in differential-algebraic equations*. PhD thesis, Ruprecht-Karls-Universität Heidelberg, 2015. Available at <http://archiv.ub.uni-heidelberg.de/volltextserver/19170/>.
- [3] Dennis Janka, Christian Kirches, Sebastian Sager, and Andreas Wächter. An SR1/BFGS SQP algorithm for nonconvex nonlinear programs with block-diagonal Hessian matrix. *submitted to Mathematical Programming Computation*, 2015.
- [4] Andreas Wächter and Lorenz T Biegler. Line search filter methods for nonlinear programming: Local convergence. *SIAM Journal on Optimization*, 16(1):32–48, 2005.
- [5] Andreas Wächter and Lorenz T Biegler. Line search filter methods for nonlinear programming: Motivation and global convergence. *SIAM Journal on Optimization*, 16(1):1–31, 2005.
- [6] Andreas Wächter and Lorenz T Biegler. On the implementation of an interior-point filter line-search algorithm for large-scale nonlinear programming. *Mathematical programming*, 106:25–57, 2006.

## A. License

This is the full license text (zlib license):

```
blockSQP -- Sequential quadratic programming for problems with
          block-diagonal Hessian matrix.
Copyright (c) 2012-2015 Dennis Janka <dennis.janka@iwr.uni-heidelberg.de>

blockSQP 2 -- Condensing, convexification strategies, scaling heuristics and more
          for blockSQP, the nonlinear programming solver by Dennis Janka.
Copyright (c) 2023-2025 Reinhold Wittmann <reinhold.wittmann@ovgu.de>
```

This software is provided 'as-is', without any express or implied warranty. In no event will the authors be held liable for any damages arising from the use of this software.

Permission is granted to anyone to use this software for any purpose, including commercial applications, and to alter it and redistribute it freely, subject to the following restrictions:

1. The origin of this software must not be misrepresented; you must not claim that you wrote the original software. If you use this software in a product, an acknowledgment in the product documentation would be appreciated but is not required.
2. Altered source versions must be plainly marked as such, and must not be misrepresented as being the original software.
3. This notice may not be removed or altered from any source distribution.